

Top 6

task-specific confidence calibrator

Bedrock deterministic pipeline

ModelHub SageMaker ModelHub SageMaker AWS

AWS foundation model Bedrock SageMaker

connector coverage live cache fixture BlockTempo planned

15 budget cache fail-closed abstain

TrustForge Hermes

20

learning event run candidate production

1. TrustForge LLM RAG OpenClaw Agent

LLM RAG Agent framework TrustForge trust boundary evidence contract run execution lineage deterministic reasoning policy candidate review activation rollback Agent decision record

```
`src/trustforge/pipeline.py` `schema.py` `execlog.py` `trust/kernel.py` `upgrade_control.py`  
`frontend/src/components/AnalysisReportView.tsx`
```

2. workspace governance console

workspace governance console incident API artifacts report evidence execution log snapshot telemetry dashboard

TrustForge

Runtime

3. analysis run

run_scope_id Ingestion claim extraction dedup corroboration Trust Kernel stance pipeline report evidence execution log Evidence Trail facts/inferences ladder contrarian could_flip technical execution journey

```
`pipeline.run()`  
`agent/orchestrator.py`  
`trust/scoring.py`  
`trust/kernel.py`  
`frontend/src/components/AnalysisReportView.tsx`
```

4. connector

OHLCV HOYA BIT/RSS Whale Alert Arkham Etherscan CoinMarketCap CoinGecko DefiLlama connector safe fetch cache reliability fetched_at degraded handling

```
`src/trustforge/ingestion/`  
`prices.py`  
`hoyabit.py`  
`news.py`  
`social.py`  
`regulatory.py`  
`onchain.py`  
`whale_trades.py`  
`cmc.py`  
`coingecko.py`  
`defillama.py`
```

5. Evidence Grouping

canonical source identity claim evidence group
URL Trust Kernel independent_sources raw document
count stance pair contrarian evidence

```
`agent/evidence_grouper.py`  
`trust/scoring.py`  
`src/trustforge_core/corroboation.py`  
`frontend/src/lib/evidenceGrouping.ts`
```

6. Bedrock

Bedrock claim extraction semantic stance
classification pipeline facts inferences evidence limits
narrative source weight Trust Kernel

Bedrock offline/sample deterministic fallback fallback live

Trust Domain

7. TrustScore

TrustScore evidence calibrated confidence
coverage evidence decision state
low-confidence abstain

8. Trust Kernel I/O

Trust Kernel frozen KernelInput/KernelOutput contract
run_kernel() IO
AWS HTTP cache provider production pipeline
scoring.py BedrockClient ingestion roadmap production
composition kernel entry point SLA

```
`docs/architecture/ARCHITECTURE-OVERVIEW.md`  
`docs/architecture/TRUST-KERNEL-BOUNDARY.md`
```

9. Asset Intrinsic Peer Metrics Whale Carbon

context plane Asset Intrinsic Peer Metrics Whale Activity Carbon token TrustScore evidence risk flag comparison operational decision

```
`asset_intrinsic.py` `peer_metrics_scheduler.py` `whale_api.py` `carbon.py`  
`frontend/src/components/AssetIntrinsicShadowPanel.tsx`
```

Learning ModelHub SageMaker

10.

calibration model Isotonic Regression logistic calibrator
candidate raw confidence coin direction feature market
outcome calibrated confidence calibration metrics Claude

11. AWS

Bedrock foundation model evidence contract source
reliability freshness half-life dedup semantics delayed ground truth
calibrator Bedrock TrustForge over-confidence provider

12.

closed-loop learning run learning event raw confidence
evidence execution metadata T+1/T+7/T+14 outcome label label
leakage training row `prepare\_calibrator\_training.py` training contract
calibrator gate backend submit candidate CLI
enhancement activation

```
`docs/reports/REPORT-2026-07-31-TRAINING-AUTO-TRIGGER-ANALYSIS.md` scheduler
```

13. [ModelHub](#) [SageMaker](#) [artifact](#) [provenance](#)

[ModelHub](#) [governance surface](#) [model version](#) [training request](#) [metrics](#) [artifact provenance](#) [candidate status](#) [approval lifecycle](#) [SageMaker AI](#) [execution backend](#) [Training Job](#) [S3](#) [artifact](#) [TrainingBackend](#) [contract](#) [ModelHub](#) [SageMaker](#) [candidate](#) [production](#)

14. [AWS foundation model](#)

[foundation model](#) [Amazon Bedrock](#) [calibrator](#) [TrustForge](#) [deterministic task model](#) [AWS](#) [SageMaker AI](#) [AWS](#) [SageMaker backend](#) [ModelHub](#)

Upgrade [AgentCore](#)

15. [Agent](#) [prompt](#) [production code](#)

[upgrade control](#) [provider](#) [policy](#) [model](#) [calibrator](#) [RAG index](#) [reranker](#) [candidate](#) [diagnostics](#) [sandbox](#) [adversarial review](#) [artifact integrity](#) [reviewer gate](#) [human activation](#) [rollback](#) [Trust Kernel](#) [reviewed code release](#) [outer skill](#) [merge](#)

[upgrade_control.py](#) [upgrade_queue.py](#) [upgrade_review.py](#) [policy/](#) [docs/architecture/UPGRADE-CONTROL-OVERVIEW.puml](#)

16. [AgentCore](#) [Memory](#) [Policy](#) [Hermes](#)

[Hermes](#) [TrustForge](#) [domain agent product](#) [AgentCore](#) [runtime/provider integration](#) [Memory plane](#) [session](#) [run](#) [evidence](#) [learning event](#) [Policy plane](#) [AgentCore](#) [TrustScore](#) [TrustForge](#) [evidence contract](#)

17. [Execution Log](#)

formal run [idempotency key](#)[lease](#)[budget reservation](#) [durable state](#)[multi-angle batch](#) [atomic batch store](#)[reconcile](#) [terminal receipt](#)[timeout](#)[remaining budget](#) [retry policy](#)[report](#) [limits](#) [data mode](#) [failed Execution Log](#)

[AWS Architecture](#)

18. [AWS Architecture](#)

[EC2](#) [t3.micro](#) [nginx](#) [80/443](#)[React/Vite](#) [nginx](#) [Python API](#) [127.0.0.1:8080](#)[Bedrock](#) [LLM](#) [S3](#) [deployment/artifact](#)[IAM Instance Role](#)[SSM Session Manager](#) [SSH key](#)[CloudWatch](#) [App Runner](#)

``docs/architecture/AWS-ARCHITECTURE.md``deploy/nginx.conf``deploy/deploy\_ec2.sh``

19. [API Key](#)[AWS credential](#)

[credential](#) [Git](#)[EC2](#) [IAM Instance Role](#)[SSM](#)[live model access](#) [token gate](#)[rate limit](#)[budget guard](#) [request accounting](#) [nginx](#) [Python API](#) [loopback](#)[admin endpoints](#) [admin token](#) [fail-closed](#) [artifact](#) [evidence](#) [secret](#)[stack trace](#)

20. [Evidence Trail](#)

[run_id](#) [claim_id](#)[Final Report](#) [Evidence List](#) [source](#)[Execution Log](#) [Source/Config](#)[artifact manifest](#) [SHA-256](#) [UI](#) [Evidence Trail](#)[Execution Journey](#)

``schema.py``execlog.py``execution_event_log.py`
`frontend/src/components/ReportDownloads.tsx``EvidenceTrailPanel.tsx``

ModelHub/SageMaker training

prompt evidence contract
Trust Kernel calibration governance
replay operational controls

ModelHub/SageMaker training

ModelHub/SageMaker training
backend
activation

Trust Layer API

Trust Layer API
outcome labels audit records

TrustForge AI

TrustForge AI
AI