

TRUSTFORGE HERMES EVIDENCE-NATIVE DECISION INTELLIGENCE

AI Agent

TrustForge Hermes evidence-native AI Agent audit trail

decision intelligence



RESEARCH WORKSPACE

Hermes Workspace

Analyze▯Compare▯History▯Source Status▯Whale Activity ▯ Evidence Trail▯▯▯▯▯▯
▯▯▯▯▯▯▯▯▯▯▯▯▯▯

OPERATIONS CONSOLE

Governance & Operations

Training Status▯Budget▯Carbon▯Module Telemetry▯Upgrade Queue▯Memory ▯
AgentCore ▯▯▯▯▯▯▯

INTEGRATION SURFACE

API & Artifacts

REST API▯comparison snapshot▯analysis journey▯Final Report▯Evidence List ▯
Execution Log ▯▯▯▯▯▯▯▯▯

Decision support—not trading

run_id / claim_id

Human-in-the-loop

Offline / Live / Degraded

Hermes Agent Runtime

Ingestion

OHLCV Whale HOYA BIT



Reasoning Pipeline

Claim extraction stance dedup corroboration Trust Kernel uncertainty



Evidence Delivery

Report Evidence Execution Log Snapshot API UI

Bounded tools

Policy plane

source analysis report evaluation improvement policy

Memory plane

Trust Kernel

Evidence contract

source[] fetched_at[] content_reference[] related_claim [] source_url []



Deterministic reasoning

source reputation[] freshness half-life[] independent corroboration[] stance [] manipulation risk[]



Decision state

direction[] trust score[] calibrated confidence[] reason codes[] abstain [] could_flip[]

Source

canonical source identity[] claim grouping [] union-find []

Stance

contrarian evidence[] Conflict Ledger [] stance pair []

Asset

Asset Intrinsic[] Peer Metrics[] multi-angle [] Whale history []



analysis run

report evidence execution log raw confidence learning event



Delayed outcomes

T+1 / T+7 / T+14 OHLCV direction label ground truth



Candidate training

Isotonic / logistic calibrator ECE holdout artifact hash proposal

ModelHub

training request metrics artifact provenance candidate lifecycle

Amazon SageMaker AI

AWS training job S3 dataset artifact AWS

run candidate automatic_apply=false



UPGRADE CONTROL

Candidate → Review → Activate

```
Provider → policy → model → calibrator → RAG/reranker → sandbox → reviewer
gate → release →
```

RELIABILITY

Idempotency & Recovery

```
formal run multi-angle batch Dynamo/SQLite lease atomic reconcile retry  
rollback receipt
```

SECURITY & COST

Zero-trust Operations

IAM least privilege □ SSM □ token gate □ rate limit □ budget reservation □
CloudWatch □ carbon ledger □ secret scan □

AgentCore runtime

Policy compiler / executor

Memory & lineage

Module telemetry

Release gates

Agent

📦 AWS 📦📦📦📦📦📦📦📦

React / Hermes UI

Vite build📦responsive workspace📦analysis📦compare📦history📦admin 📦 operational pages📦



nginx → Python API

EC2 t3.micro📦nginx :80/:443📦Python 📦 127.0.0.1:8080📦API 📦 SSR 📦📦📦📦📦 API📦



AWS & data plane

Amazon Bedrock📦S3📦IAM📦SSM📦CloudWatch📦Lambda/AgentCore 📦📦📦📦📦📦

Live / Sample / Degraded

📦 contract 📦📦📦📦📦📦📦📦📦📦 UI 📦 log 📦📦📦

API-first

health📦status📦analyze📦comparison📦training📦upgrade📦whale📦carbon📦delivery 📦
telemetry📦

Artifact-first

Report📦Evidence📦Execution Log📦Source/Config📦manifest 📦 SHA-256 📦📦📦📦📦

TrustForge

TrustForge

- decision record
- confidence limits could_flip contrarian evidence
- API artifact telemetry

TrustForge

- provider policy memory retrieval calibrator
- candidate review rollback audit
- Agent

OpenClaw Agent TrustForge evidence

TrustForge

TrustForge

4 PoC

5 4 Evidence

KPI

TrustForge

AI Agent