

TrustForge AI Agent の紹介とデモ

TrustForge Hermes の紹介とデモ

8 月 6 日 14:00-15:00

TrustForge の runtime (Trust Kernel) と training backend を
run する learning event (T+1/T+7/T+14 outcome) の candidate を
EC2 + nginx の App Runner

10:00-0:40

Product thesis の Agent の紹介とデモ

Team 11 の TrustForge の demo
TrustForge Hermes (evidence-native decision intelligence agent)

deterministic trust reasoning の audit trail
runtime の decision record

2 0:40-1:15

Product surface

Hermes Workspace Analyze Compare History Source Status
Whale Activity Evidence Trail artifacts

Training Status Budget Carbon Module Telemetry Upgrade
Queue Memory AgentCore admin audit API comparison snapshot
analysis journey Final Report Evidence List Execution Log workflow

3 1:15-2:00

Runtime architecture Agent

Hermes runtime Ingestion OHLCV Whale HOYA BIT
reasoning pipeline claim extraction stance deduplication
corroboration Trust Kernel uncertainty report evidence execution
log snapshot API response

Agent tool registry Memory plane run artifact
evidence learning event replay history policy plane source analysis
report evaluation improvement degraded abstain fail-closed

4 2:00-2:45

Trust Kernel decision state

TrustForge Evidence contract source fetched_at content_reference related_claim URL endpoint query

Trust Kernel deterministic logic freshness half-life stance manipulation risk trust score calibrated confidence direction reason codes independent source count abstain canonical source identity claim grouping union-find

contrarian evidence Conflict Ledger could_flip stance pair Asset Intrinsic Peer Metrics multi-angle Whale history

5 2:45-3:35

Learning system

confidence TrustForge Evidence contract
analysis run report evidence
execution log raw confidence learning event label leakage
T+1 T+7 T+14 OHLCV direction outcome

JSONL training contract schema finite-value split
calibrator gate Isotonic Regression logistic calibrator ECE
holdout improvement artifact hash candidate proposal

ModelHub request artifact provenance lifecycle SageMaker
AI AWS Training Job S3 dataset artifact candidate
production automatic_apply=false requires_human_approval=true

6 3:35-4:20

Control plane

Hermes upgrade control provider policy model calibrator RAG index
reranker candidate sandbox diagnostics adversarial review
artifact integrity reviewer gate activation Trust Kernel reviewed
code release outer skill

formal run multi-angle batch DynamoDB SQLite lease
atomic reconcile retry rollback receipt idempotency key
IAM least privilege SSM token gate rate limit budget reservation
CloudWatch carbon ledger secret scan release gate

7 4:20-5:05

Deployment AWS

EC2 t3.micro nginx 80/443 Python API 127.0.0.1:8080
React/Vite frontend nginx API TrustForge pipeline Amazon Bedrock
LLM S3 artifact IAM Instance Role SSM Session Manager
CloudWatch SSH key

App Runner provider storage memory
model backend EC2 + nginx

run_id Final Report Evidence List Execution Log
Source/Config manifest SHA-256

8:05-6:00

Commercial outcome

TrustForge decision record confidence limits could_flip contrarian evidence provider policy memory retrieval calibrator model upgrade candidate review rollback audit

Agent Agent evidence-native policy-aware replayable reversible

AI AI Agent