

Executive Summary

HOYA BIT Trust Layer
PoC5 Evidence → 8-12 Pilot → 3-6 Production
SLAauditbudgettenantdashboard
KPI run API usageseatenterprise integration
live/cache/fixture API EvidenceLogartifact

RAG Crypto AI BI Dashboard TrustForge
Claim → Evidence → Trust → Report

HURRICANESOFT

TrustForge Hermes

Evidence-native Decision Intelligence Agent

	HurricaneSoft		
			2026-08-01
	6		

TrustForge AI AI Agent

TrustForge Hermes evidence-native decision intelligence agent
evidence contract Claims deterministic trust

reasoning [Amazon Bedrock](#) [TrustedBrief](#) [decision record](#) [TrustForge](#) [Trust Layer](#)

TrustForge **Trust Score**

-
- claim-to-source
- API [artifact](#) [telemetry](#)
- provider [memory](#) [policy](#) [calibrator](#) [backend](#)

RAG

	ingestion claim extraction deduplication evidence grouping	Analysis Journey Final Report
	uncertainty limits could-flip	Evidence List Trust breakdown
	run_id Execution Log model version manifest	Execution Log SHA-256 manifest
	learning event outcome gate	Training contract candidate approval receipt

TrustForge [LLM prompt](#)

2.1

Ingestion Layer		pricesHOYA BITnewsocialregulatoryon-chain whale tradesCoinGeckoCMCDefiLlamaEtherscan
Reasoning Layer	uncertainty	pipelineTrust Kernelcorroborationevidence grouperfreshnessdirection resolution
Agent / Delivery Layer	API	Hermes orchestratorBedrockreport deliveryFinal ReportEvidence ListExecution LogSnapshot

2.2

Frontend React/Vite Hermes Workspace Backend Python APIanalysis flow formal run runtime API run_id

Research Workspace	AnalyzeCompareHistorySource StatusWhale ActivityEvidence Trail facts/inferences ladder report download
Operations Console	Training StatusUpgrade QueueModule TelemetryBudgetCarbonHealth
Governance Console	AdminAGOSpolicymemoryprompt versionapprovalrollbackaudit record
Integration Surface	API responsesnapshotartifactmanifestFinal ReportEvidence ListExecution LogSource/Config

2.3 AWS

EC2 t3.micro + nginx nginx HTTPS frontendPython API 127.0.0.1:8080Amazon Bedrock LLM S3 IAM Instance RoleSSM Session Manager CloudWatch App Runner

コンポーネント	機能	依存関係
EC2 + nginx	React/Vite frontend, Python API reverse proxy	TLS, 127.0.0.1 API binding
Amazon Bedrock	claim extraction, stance classification, report assembly	Trust Score
S3	artifact, manifest	IAM
IAM / SSM / CloudWatch	SSH key	

3.0 System Architecture

3.1 Data Ingestion

Ingestion: source, fetched_at, content_reference, asset, data_mode, freshness, ...

- HOYA BIT, OHLCV, delayed outcome ground truth
- RSS, API, X, Reddit
- Etherscan, Whale Alert, Arkham, CoinMarketCap
- cache, safe fetch, connector reliability, freshness, data integrity, SHA-256 manifest

3.2 Trust Kernel

Trust Kernel: deterministic trust reasoning, LLM, confidence, uncertainty

Trust Kernel, Bedrock

3.3 Formal Run

- formal_run_coordinator, budget reservation, idempotency, lease
- analysis flow: claim extraction, stance, deduplication, corroboration, Trust Kernel, uncertainty

- run run_id Final ReportEvidence ListExecution LogSnapshotSource/Config API response
- timeout degradedabstain fail-closed

3.4 Agent RuntimeMemory Policy

Hermes orchestrator AgentCore adapter/runtime tool registry memory OS policy compiler/executor Agent RuntimeMemory plane runartifactevidencelearning event replay historypolicy plane source analysis report evaluation improvement

AgentCore TrustForge domain logic TrustForge contractguards audit

3.5 Upgrade Control

Improvement upgrade control canary activation lock deployment evidence rollback record refresh_sources measure_quality replay_history diagnose_improvement candidate candidate benchmark activation

Candidate	production	metricsartifact hashtraining contract
Review	qualitysecuritycostcompatibility adversarial review	reviewer attestationdecision record
Activate	canary activation lock	approvaldeployment evidence receipt
Rollback		rollback decision recordhealth checkaudit trail

3.6

- Module Telemetry
- Budget Guard tokenBedrock Agent USD
- Carbon input/output token Trust Kernel LLM

- Audit hermes audit execution events artifact registry authenticated ledger

Bedrock

4.1

Bedrock foundation model prompt

- deterministic feature prompt
- delayed outcome confidence benchmark
- ModelHub governance SageMaker execution backend artifact
- foundation model Amazon Bedrock TrustForge task-specific calibrator foundation model

4.2 Learning Event

learning event Claims Trust Score raw confidence run_id Execution Log

outcome T+1 T+7 T+14 1714

4.3 backend

1. Run learning event immutable evidence execution metadata
2. T+1/T+7/T+14 outcome label training row label leakage
3. Isotonic Regression candidate ECE Brier holdout
4. SageMaker backend AWS S3 artifact ModelHub
5. candidate review canary activation production

training trigger CLI scheduler run activation enhancement

4.4 ModelHub SageMaker

Tool	Category	Capabilities
ModelHub	Governance surface	model version, training request, metrics, artifact provenance, candidate status, approval lifecycle
SageMaker AI	Training execution backend	via S3, orchestrate Training Job, artifact, orchestrate
S3	Artifact storage	training dataset, calibrator artifact, manifest, orchestrate

Data sources

5.1 Data sources

Category	Source	Capabilities
OHLCV	HOYA BIT, prices, CoinGecko	Retrieve T+ outcome, orchestrate
News	RSS, API, news connector	Retrieve news, orchestrate
Social	X, Reddit, social connector	Retrieve social media, orchestrate
Blockchain	Etherscan, Solscan, on-chain	Retrieve blockchain data, orchestrate
Regulatory	Taiwan regulatory	Retrieve regulatory information, orchestrate
Market Data	CMC, CoinGecko, DefiLlama	Retrieve market data, orchestrate

5.2 Evidence Contract

The Claim object contains source, fetched_at, content_reference, asset, data_mode, and related_claim. Evidence List contains a list of evidence objects, each containing source, fetched_at, content_reference, asset, data_mode, and related_claim.

5.3 Data sources

項目	内容	備考
Final Report	Trust Scoreconfidencelimitscould-flip	
Evidence List	claim_id	
Execution Log		
Source / Config + Snapshot	manifest	

- API keyAWS credentialsGitHubIAM Instance RoleSSM secret boundary
- tool registry policy guards scope
- Fail-closeddegraded abstain
- run_idmodel versionartifact hashapproval receiptdeployment evidence rollback record
- candidate adversarial reviewsecret production gate

TrustForge ISO/IEC 42001 ISO 42001

Overview

Phase	Key Components	Key Metrics
Phase 1: Research Trust Layer	Trust Layer API, report, evidence, execution, snapshot	Trust Layer API adoption, review turnaround
Phase 2: Governed Decision Services	API, peer metrics, asset context	API adoption, review turnaround
Phase 3: Continuous Calibration	outcome label, candidate, ModelHub, SageMaker, canary	ECE, Brier, candidate, rollback
Phase 4: Enterprise Control Plane	IAM, audit, budget, carbon, telemetry, policy	

Key Features

- Evidence contract, claim-to-source lineage, prompt
- Trust Kernel, backend
- Bounded autonomous agent
- Calibration, replay, artifact provenance, rollback
-

Architecture

Component	Details
API	data_mode, fixture, cache, live
API	Trust, Evidence
API	Final Report, Evidence List, Execution Log, Source/Config, manifest, run_id
API	training row, outcome label, candidate, metrics, artifact hash
API	degraded, abstain, fail-closed

- TrustForge 部署到 Kubernetes
- 在 provider 部署 AgentCore 和 DynamoDB backend 到 Kubernetes
- 在 Kubernetes 上运行 TrustForge scheduler 和 activation
- 在 Kubernetes 上部署 TrustForge 客户端

[illegible]

--	--	--	--	--	--	--

TrustForge Hermes HurricaneSoft